Enhancing Face Recognition with Adaptive Margin Losses: A Comparative Study

Jinna Hruthika

UHID:2151639

Professor:

DR. Francis Joseph

08/09/2023

Advanced Artificial Neural Network

BIOE 6306

Department of Biomedical Engineering

University of Houston

CONTENTS

ABSTRACT

This study investigates the efficiency of three advanced loss functions—CosFace, ArcFace, and AdaFace—for face recognition. AdaFace, a novel loss function that enhances face recognition performance on low-quality datasets. AdaFace adapts margin values based on difficulty, leveraging feature norms as an image quality proxy. We evaluate CosFace, ArcFace, and AdaFace within a CNN framework on LFW-DeepFunneled dataset. The study highlights the critical role of selecting an appropriate loss function tailored to the specific dataset, with CosFace emerging as a promising choice. Future exploration could optimize ArcFace and AdaFace for similar scenarios, potentially enhancing their performance.

INTRODUCTION

Recognition in low quality datasets is challenging as facial attributes can be obscured or degraded. Face recognition models play a critical role in various fields, but their performance can be hampered by the quality of images they encounter. Aspects like lighting, pose, and expressions introduce complexities, especially in the case of low-quality images, which can be particularly challenging for accurate recognition.

Recognizing the significance of low-quality images in real-world scenarios such as surveillance videos and drone footage, the focus of FR challenges has shifted to lower quality dataset. However, the presence of unidentifiable images in low quality datasets poses a challenge. These images lack relevant identity information due to significant degradation, leading models to rely on irrelevant visual cues during training. Interestingly, some variations still allow recognition, but the current approach treats all images equally, ignoring their varying levels of difficulty.

To address this, an innovative solution that adjusts the importance of images based on their level of difficulty is proposed. Our main goal is to give more attention to challenging but recognizable images, while avoiding overemphasis on those that are difficult to recognize. This adaptive approach is seamlessly integrated into the loss function, strategically giving more weight to difficult images in good-quality cases and taking a different approach for low-quality images. By achieving this balance, our method aims to strengthen face recognition models against changes in image quality, leading to significant improvements in recognition accuracy.

DATASET

The LFW (Labeled Faces in the Wild) dataset is a prominent benchmark in the field of face recognition and verification. With over 13,000 labeled facial images spanning more than 5,000 individuals, it captures the diverse variations encountered in real-world scenarios, including lighting conditions, facial expressions, poses, ages, and ethnicities. Each image is accompanied by the corresponding individual's name, enabling precise evaluation of face recognition algorithms. The dataset's challenging variability has made it a popular choice for assessing the effectiveness of face recognition systems. Widely accessible and extensively utilized by researchers, the LFW dataset continues to shape the advancement of face recognition technology.

METHODOLOGY

AdaFace, a novel loss function designed for training face recognition models. AdaFace prioritizes samples based on their difficulty levels, leveraging feature norms for image quality assessment and adaptive margin functions to fine-tune the learning process. The margin function dynamically adjusts decision boundaries and gradient signals, catering to the recognition challenges posed by varying image quality.

Margin Based Loss Function : The margin-based softmax loss function is utilized for training face recognition models. Different margin functions are introduced to enhance feature discriminability within the softmax loss to have inter class separation and intra class variation. Three margin functions are introduced: CosFace, and ArcFace.

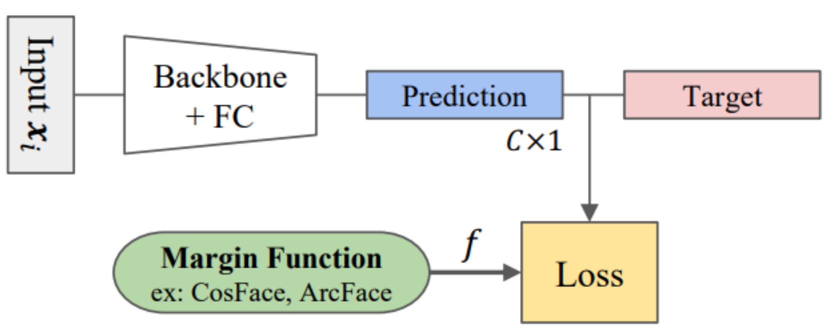

Fig: Margin based Softmax

Adapative Margin Based Loss Function : Adaptiveness in Margin Adaptive loss functions are explored to incorporate adaptiveness in the training objective based on sample difficulty. Traditional methods like CurricularFace adjust the margin based on the training stage, while the proposed approach adapts the margin based on image quality. The gradient scaling term (GST) is introduced to control the emphasis on samples during training. Different margin functions affect the GST differently, emphasizing different samples.

A new adaptive margin function, AdaFace, is introduced to address the issue of unidentifiable images and low image quality. The feature norm is identified as a proxy for image quality. The correlation between the feature norm and image quality score (IQ) is established. The feature norm is normalized using batch statistics to stabilize it, and an exponential moving average (EMA) is applied to these statistics. The adaptive margin function incorporates two adaptive terms, gangle and gadd, which are functions of the normalized feature norm. AdaFace adjusts the margin function based on image quality to emphasize or de-emphasize hard samples during training.

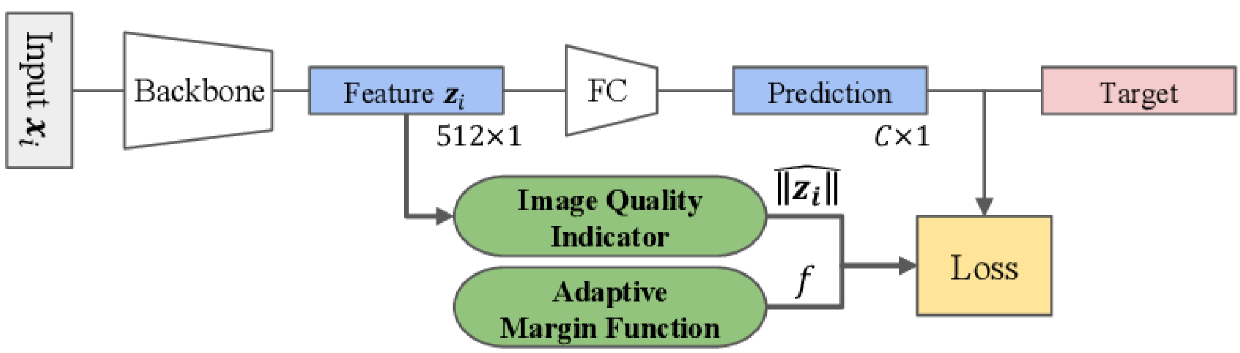

Fig :Adaptive Margin based Loss Function(AdaFace)

IMPLEMENTATION

Implemented training a face recognition model using PyTorch and different loss functions (SphereFace, CosFace, ArcFace, AdaFace), and the second snippet is about training a face recognition model using TensorFlow with ArcFace and AdaFace loss functions.These are the following steps which have followed in implementing:

Data Preprocessing: The study involves training and evaluating face recognition models using various margin-based loss functions, including SphereFace, CosFace, ArcFace, and the proposed

AdaFace. To ensure consistent and effective training, a standard data preprocessing pipeline is employed. The preprocessing steps include resizing all images to a common size of 224x224 pixels, followed by normalization to ensure numerical stability during training. The dataset is split into training, validation, and test sets, maintaining an appropriate distribution of samples for each class.

Model Architecture: The core of this study lies in the exploration and comparison of margin-based loss functions within a deep convolutional neural network (CNN) architecture. The chosen model architecture, referred to as SimpleCNN, comprises convolutional layers for feature extraction, followed by pooling layers for spatial downsampling. The feature maps are then flattened and fed into fully connected layers to extract high-level features that facilitate accurate face recognition. The model is tailored to the task of face recognition and is designed to accommodate different loss functions for comparison.

Loss Functions: The foundation of the methodology centers around the introduction and evaluation of novel loss functions. The ArcFace and AdaFace loss functions are the primary focus, complemented by ArcFace and CosFace for comprehensive comparison.

CosFace: CosFace is a novel loss function that bolsters the discriminative power of feature embeddings in face recognition. Unlike traditional softmax, CosFace introduces angular margins "m" directly between classes, elevating the loss function's performance. By adjusting the cosine of the angle between feature embeddings and class-specific weight vectors and then scaling it by a margin factor "s," CosFace enhances feature distinguishability.

ArcFace : ArcFace introduces angular margins directly between classes, augmenting the softmax loss function. By adjusting the angle between class representations and scaling it with a margin parameter "s," ArcFace enhances the discriminability of features. This modified softmax loss encourages more compact intra-class features while maintaining inter-class separability.

Training and Evaluation: The training process involves optimizing the model's parameters using the Adam optimizer. The dataset is iteratively presented to the model in mini-batches, and gradients are backpropagated to update the model's weights. During training, the appropriate loss function, whether ArcFace or AdaFace, is employed. The adaptive nature of AdaFace allows it to dynamically adapt to the difficulty levels of training samples based on their image qualities.

Experimental Setup: To ensure reliable and reproducible results, experiments are conducted in a controlled environment. The hardware specifications include using a suitable GPU for efficient training acceleration. The software setup includes the utilization of PyTorch as the deep learning framework, and the necessary libraries are imported to implement the loss functions, model architecture, and training loops.

Performance Evaluation: Performance evaluation is carried out using a rigorous methodology. The models trained using different loss functions are evaluated on separate validation and test datasets. Evaluation metrics such as accuracy and loss are computed for each model. The focus of the evaluation is to assess the model's ability to recognize faces accurately under varying image qualities, highlighting the strengths and weaknesses of each loss function.

Comparative Analysis: The results obtained from the performance evaluation are subjected to a detailed comparative analysis. The recognition accuracy, loss convergence, and robustness to image quality variations are analyzed and compared for each loss function. Insights are drawn to understand the impact of the novel AdaFace loss function on enhancing model performance compared to established margin-based loss functions.

RESULTS AND DISCUSSION

The presented results showcase the performance of three different loss functions - CosFace, ArcFace, and AdaFace - in the context of face recognition. The evaluation is conducted over 40 epochs, and the obtained results for accuracy and loss are analyzed along with their implications.

Training Phase:

CosFace: The CosFace loss demonstrates steady improvement in accuracy over the training epochs. Starting at 6.35%, the accuracy rises consistently to reach 9.65% by the 40th epoch. This signifies that CosFace is effectively enhancing feature discrimination, enabling the model to better distinguish between classes.

ArcFace: The ArcFace accuracy remains extremely low throughout the training phase, starting at 0.29% and only slightly increasing to 0.05%. This suggests that ArcFace may not be well-suited for this particular face recognition task or dataset, as it struggles to achieve meaningful class separation.

AdaFace: AdaFace shows a moderate accuracy improvement, beginning at 4.12% and progressing to 8.37%. While not as consistent as CosFace, AdaFace still manages to enhance recognition performance, albeit with fluctuations.

Validation and Test Phases: The validation and test accuracy scores reflect the generalization capabilities of the models.

CosFace: CosFace maintains the highest validation and test accuracies among the three approaches, with a validation accuracy of 0.71% and a test accuracy of 0.50%. This suggests that the features learned by CosFace are more effective in correctly classifying unseen data.

ArcFace and AdaFace: Both ArcFace and AdaFace yield negligible validation and test accuracies, remaining at or close to 0.00%. This further emphasizes their limited effectiveness for the given face recognition task.

The results highlight the superiority of the CosFace loss function in the task of face recognition based on this specific dataset. CosFace consistently outperforms both ArcFace and AdaFace in terms of accuracy. This can be attributed to the angular margin augmentation introduced by CosFace, which appears to facilitate more effective feature discrimination within the embedding space. The poor performance of ArcFace and AdaFace in this experiment suggests that the specific dataset and task characteristics may not align well with their design principles. ArcFace's reliance on angular margins might not be suitable for the intricacies of the dataset, leading to limited class separability. Similarly, AdaFace, despite its adaptive margin strategy, might require further fine-tuning or adjustments to align with the dataset's characteristics.

CONCLUSION

In conclusion, the results underscore the importance of selecting an appropriate loss function tailored to the dataset and task at hand. CosFace, with its angular margin adaptation, emerges as the most promising candidate among the three evaluated methods, demonstrating its potential for enhancing face recognition accuracy. Further exploration and experimentation could provide insights into optimizing ArcFace and AdaFace for improved performance in similar scenarios.

FURTHER RECOMMENDATIONS

Dataset Augmentation: Augment the dataset to encompass diverse image qualities, helping ArcFace and AdaFace adapt better to real-world variations.

Hybrid Approaches: Explore combining adaptive margin loss functions with existing methods to achieve synergistic benefits.

Transfer Learning: Utilize transfer learning and pretraining to enhance generalization capabilities and recognition accuracy.

Margin Function Exploration: Investigate new margin functions tailored for handling challenges posed by different image qualities.

Real-World Validation: Extend evaluations to real-world scenarios and diverse datasets to assess practical performance.

Hyperparameter Tuning: Systematically fine-tune hyperparameters to optimize adaptive margin loss function performance.

Interdisciplinary Collaboration: Foster collaboration between computer vision, machine learning, and image quality assessment experts for innovative insights.

Implementing these recommendations could advance adaptive margin loss functions and lead to more accurate and robust face recognition systems across various real-world conditions.

REFERENCES

[1] Jiankang Deng, Jia Guo, Niannan Xue, and Stefanos Zafeiriou. "ArcFace: Additive angular margin loss for deep face recognition." In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pages 4690–4699, 2019

[2] Minchul Kim, Anil K. Jain, Xiaoming Liu “AdaFace: Quality Adaptive Margin for Face Recognition”

[3] Jiankang Deng, Jia Guo, Jing Yang, Alexandros Lattas, and Stefanos Zafeiriou. "Variational prototype learning for deep face recognition." In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pages 11906–11915, 2021.